

ArabicaAI intake: status, the AI decision, and who does what

30 Sep 2026 · Lucas Tramonte, Roberto, Manoella · For Roberto now, and for Factored's review tomorrow · Submission due 5 Oct

WHERE WE ARE

- **Data.** S3 → Bronze → Silver → quality gate → Gold slice, all in DuckDB. The gate passes on the full build, and each of the 18 registered findings has its own query.
- **Service.** One Cloudflare Worker with D1. Identity comes only from the session, retrieval is scoped to the customer, and the case is idempotent. The reference is issued only after the saved case is read back. Agents get a read-only queue. Adversarial tests cover the gates, sessions, isolation, idempotency and D1 budgets.
- **Demo data.** Still fictitious charges, plus one transaction from the dataset.
- **Evaluation.** A checklist baseline, an episode scorer, and a blind frozen ES/PT set of 60 cases. We committed its hashes before any system saw it.
- **AI.** Extractor v1 (Workers AI, gpt-oss-20b) reads the message, and the deterministic policy decides. On the development split it got 16/18 correct, 0 unsafe and 48/48 schema-valid, at about \$0.0005 per call. Its p95 is 3.25 s over all 48 calls, which trips our 3 s trigger. The checklist gets 15/25 on the authored cases. It misses currency words, non-ISO and relative dates, and paraphrases.

TARGET WORKFLOW AND WHAT'S MISSING

Validate session → retrieve permitted candidates → classify the outcome (one or several, none or incomplete, tool error) → confirm or clarify → idempotent case with evidence and pending questions → verify, then issue the reference → human review. The offline evaluation scores every attempt, not just the cases that got saved.

Built: session, scoped retrieval, idempotent case, reference after read-back, agent queue.

Missing: outcome classification; technical and incomplete handoff cases; a bounded server retry with a safe fallback; agent case detail and status history; the ES/PT interface; and service events, so episodes are scored over all attempts.

Out of scope, by design: the system never refunds, blocks a card or decides fraud. A handoff is not a resolution.

DECISION: WHAT SHOULD AI DO IN THE ONLINE FLOW, AND WHAT EARNS IT A PLACE?		
Option	What it means	Trade-off
A · Deterministic only	The checklist reads every message.	The simplest and fastest option. It fails the brief's learned-component requirement and keeps the checklist's 10/25 misses.
B · AI extracts, policy decides recommended (ADR-006)	The model turns text into facts. It never grants permissions, picks a charge or writes to the case.	Easy to audit. We can compare it with the checklist on the same cases, with the same policy. It adds one model call of latency.
C · AI runs the conversation	The model also asks the clarifying questions and writes the replies.	The most natural for the customer. It also has the most latency, cost and safety surface, and it's much harder to measure against a baseline.
Ship criteria, fixed before the frozen run: beat the checklist on the frozen set (McNemar, Wilson intervals); no increase in unsafe outcomes; p95 within the agreed budget; cost per episode reported. If AI misses any of these, the demo runs the checklist and we publish the result anyway. Roberto decides the shape and implements it.		

WHO DOES WHAT

- **Roberto.**
 1. Redesign the Angular front end, both the customer and agent views, to the standard of your Platzi-winning project. It needs ES/PT, clear outcome states, the case reference and the agent's case view.
 2. Close the workflow gaps above in the Worker.
 3. Decide and build the online AI flow: where the call sits, the timeout and checklist fallback, what gets logged, and how it shows in the UI.

One limit: you've seen frozen cases, so the extractor's prompt and parsing stay with the blind build, with Manoella's approval (ADR-006).
- **Manoella.** Bring the supplied dataset from S3 into the Worker as a reviewed Gold cohort in D1, so the demo stops using fictitious charges. She also decides the latency, development-label and escalation questions on #26.
- **Lucas.** Integrate everything, run the pre-registered frozen evaluation once per system, and assemble the submission.

QUESTIONS FOR FACTORED

1. Is unrecognized-charge intake with a human handoff the right V1 scope?
2. Does "the model extracts, the policy decides", measured against a checklist on a blind frozen set, meet the bar for the learned component?
3. Can we show real data with a demo cohort of about 1,000 customers from your dataset, served from D1?
4. What latency would a bank accept for this step? Our current budget is 3 s at p95.